

# ITSM Analytics Pipeline

This project implements a complete data pipeline for analyzing IT Service Management (ITSM) ticket data from ServiceNow. The pipeline includes data ingestion, transformation, and visualization using a modern data stack: Apache Airflow, DBT, PostgreSQL, and Apache Superset.

## Architecture

The pipeline follows a typical modern data architecture pattern:

1. **Data Ingestion:** Raw CSV data is loaded into PostgreSQL using an Airflow DAG
2. **Data Transformation:** DBT models clean, transform, and aggregate the data
3. **Data Visualization:** Apache Superset connects to the transformed data to provide analytics dashboards

## Project Structure

 Copy

```
itsm-analytics/  
├── airflow/  
│   └── dags/  
│       └── itsm_data_pipeline.py  
├── dbt/  
│   ├── models/  
│   │   ├── staging/  
│   │   │   └── stg_tickets.sql  
│   │   └── marts/  
│   │       ├── resolution_time_by_category.sql  
│   │       ├── closure_rate_by_group.sql  
│   │       ├── monthly_ticket_summary.sql  
│   │       └── open_tickets_by_priority.sql  
│   ├── dbt_project.yml  
│   └── schema.yml  
├── superset/  
│   └── dashboards/  
│       └── itsm_performance_dashboard.json  
└── README.md
```

## Setup and Installation

### Prerequisites

- Docker and Docker Compose
- Git

- Python 3.8+

## Step 1: Clone the Repository

bash

 Copy

```
git clone https://github.com/your-username/itsm-analytics.git
cd itsm-analytics
```

## Step 2: Configure Environment

Create a `.env` file with the necessary configurations:

 Copy

```
POSTGRES_USER=postgres
POSTGRES_PASSWORD=postgres
POSTGRES_DB=itsm
AIRFLOW_UID=50000
```

## Step 3: Start the Services

bash

 Copy

```
docker-compose up -d
```

This will start PostgreSQL, Airflow, and Superset.

## Step 4: Initialize DBT

bash

 Copy

```
cd dbt
dbt deps
dbt seed --profiles-dir .
```

## Step 5: Configure Airflow Connections

1. Navigate to <http://localhost:8080> (default credentials: airflow/airflow)
2. Go to Admin -> Connections
3. Add a new PostgreSQL connection:
  - Conn Id: postgres\_default
  - Conn Type: Postgres
  - Host: postgres

- Schema: itsm
- Login: postgres
- Password: postgres
- Port: 5432

## Step 6: Configure Airflow Variables

In the Airflow UI, go to Admin -> Variables and add:

- Key: itsm\_data\_path Value: /opt/airflow/data
- Key: dbt\_project\_path Value: /opt/airflow/dbt

## Step 7: Import Superset Dashboard

1. Navigate to <http://localhost:8088> (default credentials: admin/admin)
2. Go to Dashboards -> Import Dashboards
3. Upload the `superset/dashboards/itsm_performance_dashboard.json` file

## Running the Pipeline

### Manual Execution

To manually trigger the pipeline:

1. In the Airflow UI, navigate to DAGs
2. Find the `itsm_data_pipeline` DAG
3. Click the "Play" button to trigger it

### Scheduled Execution

The DAG is configured to run daily at 1:00 AM. You can modify the schedule in the DAG file.

## DBT Models

### Staging Layer

- `stg_tickets`: Cleans and standardizes the raw ticket data

### Marts Layer

- `resolution_time_by_category`: Calculates average resolution time by category and priority
- `closure_rate_by_group`: Calculates the ticket closure rate by assigned group
- `monthly_ticket_summary`: Creates monthly aggregations of ticket metrics
- `open_tickets_by_priority`: Provides a view of currently open tickets by priority

# Dashboard

The Superset dashboard includes:

1. **Ticket Volume Trends:** Line chart showing daily ticket creation
2. **Resolution Time:** Bar chart comparing average resolution time across Categories
3. **Closure Rate:** Pie chart of closure rate by Assigned Group
4. **Ticket Backlog:** Table displaying open tickets grouped by Priority
5. **Filters:** Week, Category, and Priority filters

## Assumptions

1. The CSV data follows a consistent format with the specified columns
2. PostgreSQL is used as the data warehouse
3. Tickets have a linear lifecycle with defined status values
4. The "Resolved Date" and "Resolution Time" fields may contain null values for open tickets